

MPI

- Message Passing Interface
- standard library interface specified by MPI forum
 - well recognized - each cluster has support for MPI
 - all operations include routine calls
- implements message passing model
 - data transfer
 - synchronization
- support
 - official support for C/C++, Fortran
 - available other bindings (Java, Python) but are not standard
- versions
 - MPI-1.2 (mostly sufficient)
 - MPI-2.1 (I/O, dynamic process management)
 - MPI-3 (enhanced collectives, multithreaded programming, performance tools)
 - MPI-4 (persistent send/receive)
 - MPI-5 (interoperability layer for compatibility between different implementations)
 - implementations
 - MPICH
 - OpenMPI

MPI Advantages

- recognized and standardized library, some implementations are free
- well understood
- tuned for all sorts of hardware
- used in many applications
- transferrable on code-level
- it supports many useful functions
- quite simple to use


Programming Concepts

- usually, we prepare one parallel program, which consists of many parallel processes; differentiation is made inside the code

- each parallel process has its own address space

- data is sent explicitly among processes

- a programmer manages data distribution
- a programmer takes care of data transfers



- inter-process communication consists of

- data transfers
 - point-to-point
 - collectives
- synchronization

- process

- processes reside in separate address spaces
- many processes can run on a single core or a single node
- process has its own
 - program code and program counter
 - stack and stack pointer
 - heap
 - static variables
 - may have one or multiple threads

- communicator

- processes are collected into groups (a set of processes)
- each message is sent in a context and must be received in the same context (same channel ID)
- a group and a context form a communicator
 - default communicator MPI_COMM_WORLD contains all processes
 - a process inside a communicator is identified by its rank
 - a rank is a number in the interval [0, P-1], where P is the number of processes

- message

- a message consists of
 - communicator
 - source address
 - destination address
 - tag
 - data
- a message is accompanied by a user-defined integer tag to assist the receiving process in identifying the message
 - communicator, source, destination, and tag must match
 - MPI_ANY_TAG does not care about the tag in a receiving message
 - MPI_ANY_SOURCE does not care about the source of a receiving message
- data in a message is described with MPI data types
 - compatibility
 - it is not necessary that sending and receiving datatype match

Environment

- each process must see executable and data
 - cluster middleware (SLURM) and network file system
 - processes must have permission to run over a network

- on Arnes cluster use module OpenMPI

- executable mpirun (mpirun) starts the requested number of processes

- a common approach is to use one executable on all nodes, differentiation is made inside the code

- in C/C++ we must include library mpi.h

- we compile with mpicc

- each function returns an error code or MPI_SUCCESS

- by default, an error causes all processes to abort

MPI Programming Basics

- Initialization
 - MPI_Init(&argc, &argv)
 - MPI_Init_thread is recommended with MPI-2
 - MPI_THREAD_SINGLE has the same behavior as MPI_Init
 - additional arguments that take care of thread safety needed to use OpenMP with MPI
 - it must be the first MPI routine in a program
- Finalization
 - MPI_Finalize()
 - it must be the last MPI routine in a program
- communicator enquiry
 - number of processes in communicator:
MPI_Comm_size(MPI_COMM_WORLD, &size)
 - rank of the process: MPI_Comm_rank(MPI_COMM_WORLD, &myid)
- timing
 - wall-clock time in seconds: MPI_Wtime()
 - difference between two points in time is only relevant
 - timer resolution: MPI_Wtick()

MPI Programming: Point-to-point Communication

- approaches
 - cooperative
 - data is explicitly sent by one process and received by another
 - any change in the receiver's memory is made with his participation
 - communication and synchronization are combined
 - MPI_Send/MPI_Recv
 - commonly used
 - one-sided
 - only one process takes care of data transfer
 - remote memory reads and writes
 - communication and synchronization are decoupled
 - MPI_Put/MPI_Get
 - rarely used
 - we will focus on cooperative approach


Blocking Communication

- data transfer

- MPI_Send copies data to local buffer sending process continues executing
- MPI_Recv of the receiving process gets data when data is ready
 - blocks program execution until data is ready
 - details about the received message can be obtained from the returned status (MPI_Status)

- buffer mode

- dead-lock prevention
- smoother operation
- buffer mode sending can be explicitly specified with MPI_Bsend
- buffer is of limited size; if a message is larger than the buffer, synchronous (buffer-less) mode is used



- synchronous (buffer-less) mode

- it is better to avoid copying to buffer
- sending process and receiving process must synchronously exchange data - sending process is blocked until receiving process is not ready
- synchronous sending can be explicitly specified with MPI_Ssend



- coding explicit transfers of data

- approach 1:

Process 0	Process 1
MPI_Send() to Process 1	MPI_Send() to Process 0
MPI_Recv() from Process 1	MPI_Recv() from Process 0

- works in buffered mode
- fails in synchronous mode
- unsafe as it depends on the buffer size

- approach 2

Process 0	Process 1
MPI_Send() to Process 1	MPI_Recv() from Process 0
MPI_Recv() from Process 1	MPI_Send() to Process 0

- each send has corresponding receive
- works in both modes

- example: [hello.c](#), [hello.sh](#)

- compiling on the Arnes cluster
 - module load OpenMPI
 - mpicc -o hello hello.c
- running
 - sbatch hello.sh

- examples:

- [p2p.sh](#): SLURM script for running examples
- [p2p1.c](#): blocking communication, approach 1
- [p2p2.c](#): blocking communication, approach 2
- [p2p5.c](#): blocking communication, send and receive combined in one function

Immediate or Non-blocking Communication

- immediate function calls

- function returns immediately, send is performed by another thread
- on return function gives handler for testing on completion

- it is safe to use

- not necessarily faster
- lots of short messages, lots of handlers

- not necessarily concurrent/asynchronous

- depends on MPI implementation

- function set

- request:
 - MPI_Request
- send/receive (one of arguments is request)
 - MPI_Isend, MPI_Irecv, MPI_Ibsend, MPI_Issend
- inquiry about request
 - MPI_Wait, MPI_Test

- blocking and immediate functions can be combined

- examples:

- [p2p.sh](#): SLURM script for running examples
- [p2p3.c](#): immediate communication with MPI_Wait
- [p2p4.c](#): immediate communication with MPI_Test